

Lecture 22: Support vector machines

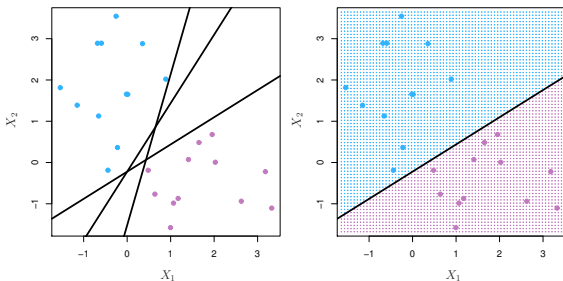
Reading: Sections 9.3, 9.4

STATS 202: Data mining and analysis

Sergio Bacallado
November 13, 2013

Maximal margin classifier

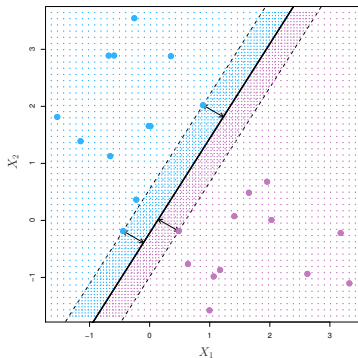
- ▶ Suppose we have a classification problem with response $Y = -1$ or $Y = 1$.
- ▶ If the classes can be separated, most likely, there will be an infinite number of hyperplanes separating the classes.



Maximal margin classifier

Idea:

- ▶ Draw the largest possible empty margin around the hyperplane.
- ▶ Out of all possible hyperplanes that separates the 2 classes, choose the one with the widest margin.

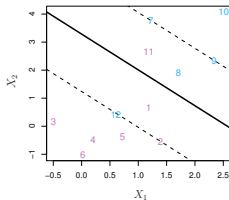
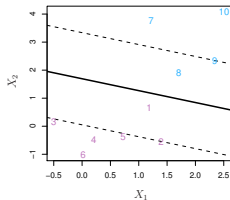


Support vector classifier

Problem: It is not always possible to separate the points using a hyperplane.

Support vector classifier:

- ▶ Relaxation of the maximal margin classifier.
- ▶ Allows a number of points to be on the wrong side of the margin or the margin or even the hyperplane.



Support vector classifier

This can be written as an optimization problem:

$$\max_{\beta_0, \beta, \epsilon} M$$

$$\text{subject to } \sum_{j=1}^p \beta_j^2 = 1,$$

$$\underbrace{y_i(\beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip})}_{\text{How far is } x_i \text{ from the hyperplane}} \geq M(1 - \epsilon_i) \quad \text{for all } i = 1, \dots, n$$

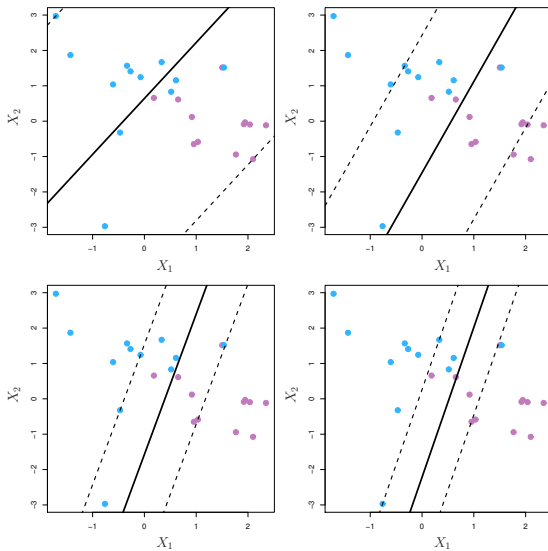
$$\epsilon_i \geq 0 \text{ for all } i = 1, \dots, n, \quad \sum_{i=1}^n \epsilon_i \leq C.$$

M is the width of the margin in either direction.

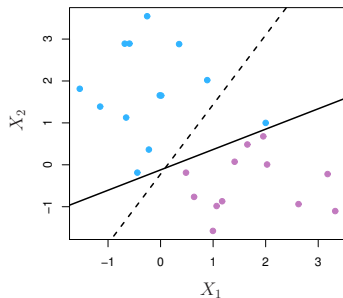
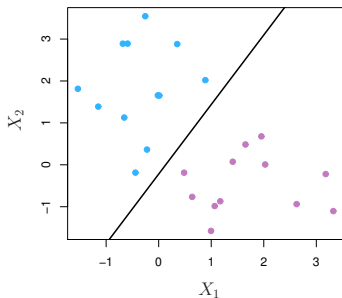
$\epsilon = (\epsilon_1, \dots, \epsilon_n)$ are called *slack* variables.

C is called the *budget*.

Tuning the budget, C (high to low)



If the budget is too low, we tend to overfit



Maximal margin classifier, $C = 0$. Adding one observation dramatically changes the classifier.

Finding the support vector classifier

We can reformulate the problem by defining a vector

$$w = (w_1, \dots, w_p) = \beta/M:$$

$$\min_{\beta_0, w, \epsilon} \quad \frac{1}{2} \|w\|^2 + D \sum_{i=1}^n \epsilon_i$$

subject to

$$y_i(\beta_0 + w \cdot x_i) \geq (1 - \epsilon_i) \quad \text{for all } i = 1, \dots, n,$$

$$\epsilon_i \geq 0 \quad \text{for all } i = 1, \dots, n.$$

The penalty $D \geq 0$ serves a function similar to the budget C , but is inversely related to it.

Finding the support vector classifier

$$\min_{\beta_0, w, \epsilon} \quad \frac{1}{2} \|w\|^2 + D \sum_{i=1}^n \epsilon_i$$

subject to

$$y_i(\beta_0 + w \cdot x_i) \geq (1 - \epsilon_i) \quad \text{for all } i = 1, \dots, n.$$

$$\epsilon_i \geq 0 \quad \text{for all } i = 1, \dots, n.$$

Introducing Lagrange multipliers, α_i and μ_i , this is equivalent to:

$$\max_{\alpha, \mu} \min_{\beta_0, w, \epsilon} \quad \frac{1}{2} \|w\|^2 - \sum_{i=1}^n \alpha_i [y_i(\beta_0 + w \cdot x_i) - 1 + \epsilon_i] + \sum_{i=1}^n (D - \mu_i) \epsilon_i$$

subject to $\alpha_i \geq 0, \mu_i \geq 0, \quad \text{for all } i = 1, \dots, n.$

Finding the support vector classifier

$$\max_{\alpha, \mu} \min_{\beta_0, w, \epsilon} \frac{1}{2} \|w\|^2 - \sum_{i=1}^n \alpha_i [y_i(\beta_0 + w \cdot x_i) - 1 + \epsilon_i] + \sum_{i=1}^n (D - \mu_i) \epsilon_i$$

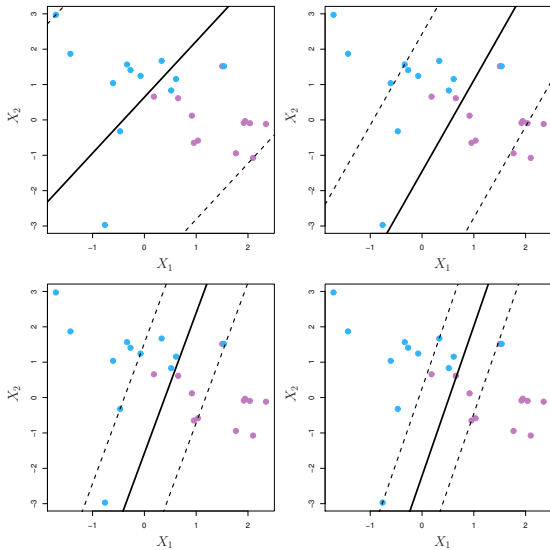
subject to $\alpha_i \geq 0, \mu_i \geq 0$, for all $i = 1, \dots, n$.

Setting the derivatives with respect to w to 0, we obtain that the solution is of the form:

$$\hat{w} = \sum_{i=1}^n \alpha_i y_i x_i$$

Furthermore, $\alpha_i > 0$ if and only if $y_i(\beta_0 + w \cdot x_i) \leq 1$, that is, if x_i falls on the wrong side of the margin.

Support vectors



The problem only depends on $x_i \cdot x_{i'}$

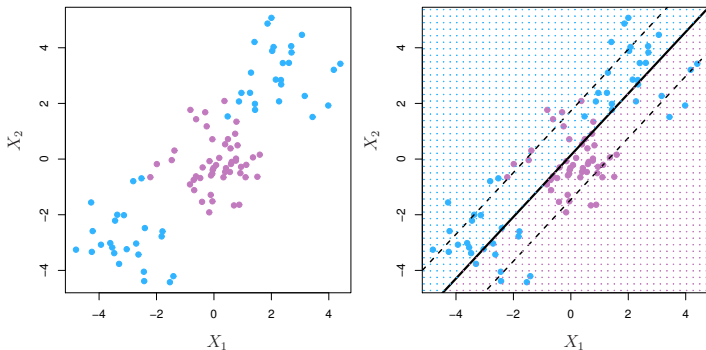
As with the Maximal Margin Classifier, the problem can be reduced to finding $\alpha_1, \dots, \alpha_n$:

$$\begin{aligned} \max_{\alpha} \quad & \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{i'=1}^n \alpha_i \alpha_{i'} y_i y_{i'} (x_i \cdot x_{i'}) \\ \text{subject to} \quad & 0 \leq \alpha_i \leq D \text{ for all } i = 1, \dots, n, \\ & \sum_i \alpha_i y_i = 0. \end{aligned}$$

As before, this only depends on the training sample inputs through the inner products $x_i \cdot x_j$ for every pair i, j .

How to deal with non-linear boundaries?

The support vector classifier can only produce a linear boundary.



How to deal with non-linear boundaries?

- ▶ In **logistic regression**, we dealt with this problem by adding transformations of the predictors.
- ▶ The original decision boundary is a line:

$$\log \left[\frac{P(Y = 1|X)}{P(Y = 0|X)} \right] = \beta_0 + \beta_1 X_1 + \beta_2 X_2 = 0.$$

- ▶ With a quadratic predictor, we get a quadratic boundary:

$$\log \left[\frac{P(Y = 1|X)}{P(Y = 0|X)} \right] = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1^2 = 0.$$

How to deal with non-linear boundaries?

- ▶ With a **support vector classifier** we can apply a similar trick.
- ▶ The original decision boundary is the hyperplane defined by:

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 = 0.$$

- ▶ If we expand the set of predictors to the 3D space (X_1, X_2, X_1^2) , the decision boundary becomes:

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1^2 = 0.$$

- ▶ This is in fact a linear boundary in the 3D space. However, we can classify a point knowing just (X_1, X_2) . The boundary in this projection is quadratic in X_1 .

How do we expand the space of predictors?

- ▶ **Idea:** Add polynomial terms up to degree d :

$$Z = (X_1, X_1^2, \dots, X_1^d, X_2, X_2^2, \dots, X_2^d, \dots, X_p, X_p^2, \dots, X_p^d).$$

- ▶ Does this make the computation more expensive?
- ▶ Recall that all we need to compute is the dot product:

$$x_i \cdot x_k = \langle x_i, x_k \rangle = \sum_{j=1}^p x_{ij} x_{kj}.$$

- ▶ With the expanded set of predictors, we need:

$$z_i \cdot z_k = \langle z_i, z_k \rangle = \sum_{j=1}^p \sum_{\ell=1}^d x_{ij}^{\ell} x_{kj}^{\ell}.$$

Kernels

A **kernel** is a matrix defined by $K(i, k) = \langle z_i, z_k \rangle$, for a set of linearly independent vectors z_1, \dots, z_n .

K is a **positive definite** matrix, i.e. it is symmetric and has positive eigenvalues.

Theorem:

If K is a positive definite $n \times n$ matrix, there exist vectors (z_1, \dots, z_n) in some space \mathbf{Z} , such that $K(i, k) = \langle z_i, z_k \rangle$.

The kernel trick

Expand the set of predictors:

- ▶ Find a mapping Φ which expands the original set of predictors X_1, \dots, X_p . For example,

$$\Phi(X) = (X_1, X_2, X_1^2)$$

- ▶ For each pair of samples, compute:

$$K(i, k) = \langle \Phi(x_i), \Phi(x_k) \rangle.$$

Define a kernel:

- ▶ Prove that a function $f(\cdot, \cdot)$ is positive definite. For example:

$$f(x_i, x_k) = (1 + \langle x_i, x_k \rangle)^2.$$

- ▶ For each pair of samples, compute:

$$K(i, k) = f(x_i, x_k).$$

- ▶ Often much easier!

The kernel trick

Example. The polynomial kernel with $d = 2$:

$$K(x_i, x_k) = f(x_i, x_k) = (1 + \langle x_i, x_k \rangle)^2$$

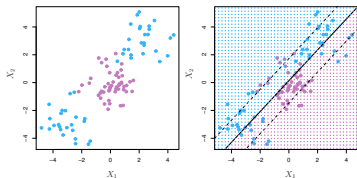
This is equivalent to the expansion:

$$\Phi(X) = (X_1, \dots, X_p, X_1^2, \dots, X_p^2, X_1X_2, X_1X_3, \dots, X_{p-1}X_p)$$

- ▶ Computing $K(x_i, x_k)$ directly is $O(p)$.
- ▶ Computing the kernel using the expansion is $O(p^2)$.

How are kernels defined?

- ▶ Proving that a bilinear function $f(\cdot, \cdot)$ is positive definite (PD) is not always easy.
- ▶ However, we can easily define PD kernels by combining those we are familiar with:
 - ▶ Sums and products of PD kernels are PD.
- ▶ Intuitively, a kernel $K(x_i, x_k)$ defines a *similarity* between the samples x_i and x_k . This intuition can guide our choice in different problems.



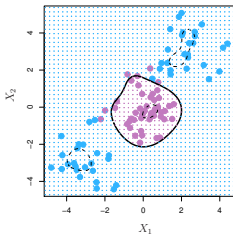
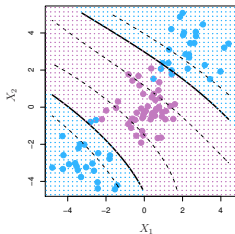
Common kernels

- ▶ The polynomial kernel:

$$K(x_i, x_k) = (1 + \langle x_i, x_k \rangle)^d$$

- ▶ The radial basis kernel:

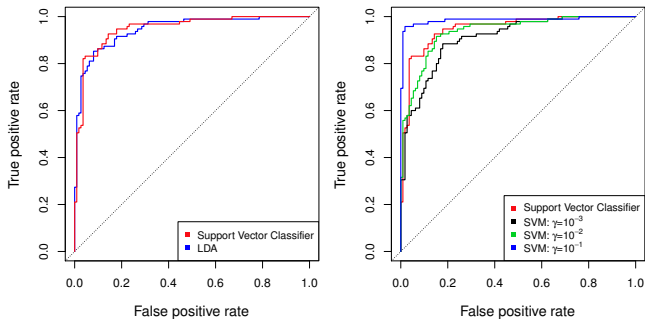
$$K(x_i, x_k) = \exp \left(- \gamma \underbrace{\sum_{j=1}^p (x_{ip} - x_{kp})^2}_{\text{Euclidean } d(x_i, x_k)} \right)$$



Dealing with non-standard data types

- ▶ Kernels are particularly useful for dealing with data types that cannot easily be represented as vectors, such as:
 1. Strings (gene sequences, search queries)
 2. Graphs (social networks)
 3. Trees
 4. Images, videos.
- ▶ It is easier to define a similarity measure which is PD than a set of features that capture the information in each sample.

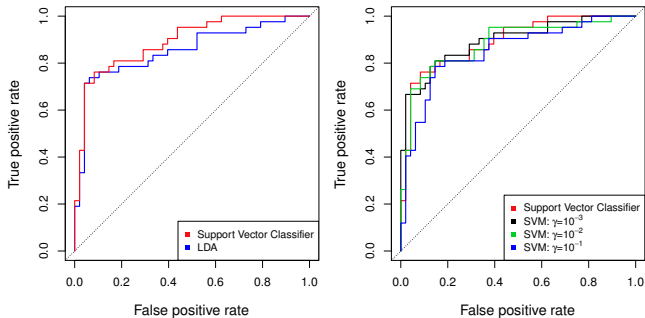
Example. Heart disease dataset



ROC curves computed on the training set.

The SVM uses a radial basis function kernel with different γ 's.

Example. Heart disease dataset



ROC curves computed on the test set.